

ISTIO

A service mesh for microservices at scale

Yong Feng – ICP STSM
yongfeng@ca.ibm.com



AGENDA

1. Microservices Architecture

- The Problem Space & Challenges


2. ISTIO

- What Is It?
- Architecture Overview

3. Demo

MICROSERVICES ARCHITECTURE

An **engineering approach** focused on decomposing an application into **single function** modules with **well defined interfaces** which are **independently** deployed and operated by **a small team** who owns the **entire lifecycle** of the service.



THE TRADE OFF

Improved delivery **velocity** in exchange for increased operational **complexity**.

Containers and Kubernetes are great enablers to these **design goals**: clean packaging, rapid deployment, consistency, reliability & scalability

This is the **reality** of microservices implementation **at scale**.

Kubernetes and containers in of themselves does not address these **complexity** challenges



Hailo microservices

Cited from <https://medium.com/@mattheath/a-long-journey-into-a-microservice-world-a714992d2841>

MICROSERVICES ARE HARD

- Applications aren't running in green-field environments
- Challenges in the network in between the Services
- Network layer is hard to manage
- Tooling is nascent

Things to consider

- Security
- Canary deployments
- A/B testing
- Circuit breaking
- Rate limiting
- Fault injection
- Tracing
- Monitoring
- Many more....

It's doable, but...

It will require a lot of coding

Service Mesh

A dedicated infrastructure layer for managing service-to-service communication to make it manageable, visible and controlled

A Control Plane & Network Overlay in between the Services

Istio

A service mesh designed to connect, manage and secure micro services.

Using **Open Source & Open Standards**

(Joint project between IBM, Google, Lyft & others)

With **Zero Application Code Changes**

MAIN FEATURES



Traffic management

A/B tests, Canary Releases, Red/Black deployments, Circuit Breaker, Fault Injection



Observability

Dependancies and traffic, Distributed Tracing, Performance metrics



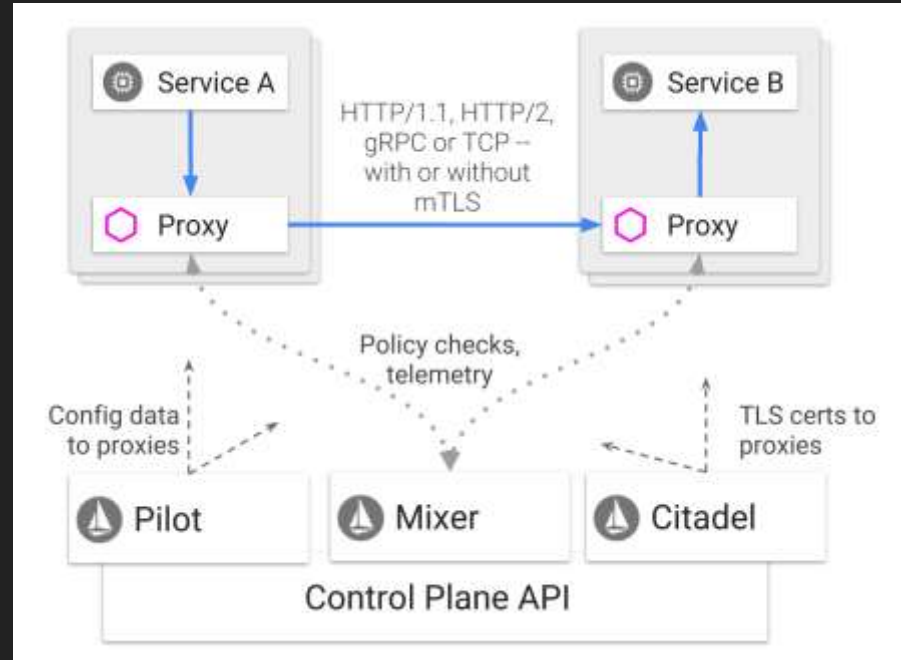
Security

Enterprise access policy, Security policy, Certification management

ISTIO CONCEPTS

Data Plane: Composed of a set of intelligent proxies (Envoy) deployed as sidecars. These proxies mediate and control all network communication between microservices along with Mixer, a general-purpose policy and telemetry hub.

Control Plane: Manages and configures the proxies to route traffic. Additionally, the control plane configures Mixers to enforce policies and collect telemetry.



Proxy: Based on Envoy, mediates inbound and outbound traffic for all Istio-managed services.

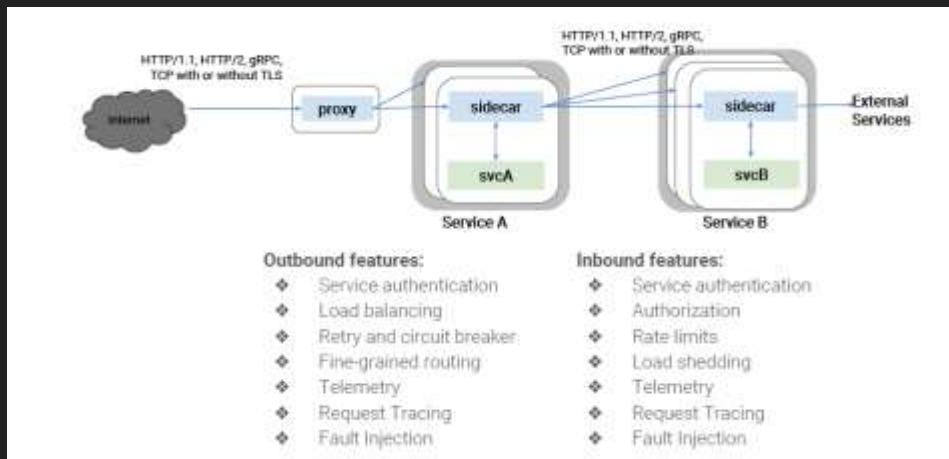
Pilot: Configures Istio deployments and propagate configuration to the other components of the system.

Mixer: Responsible for policy decisions and aggregating telemetry data from the other components in the system using a flexible plugin architecture.

Citadel: Secures the service-to-service communication and provides a key management system to manage keys and certificates.

ENVOY

SIDE CAR TECHNOLOGY OF CHOICE



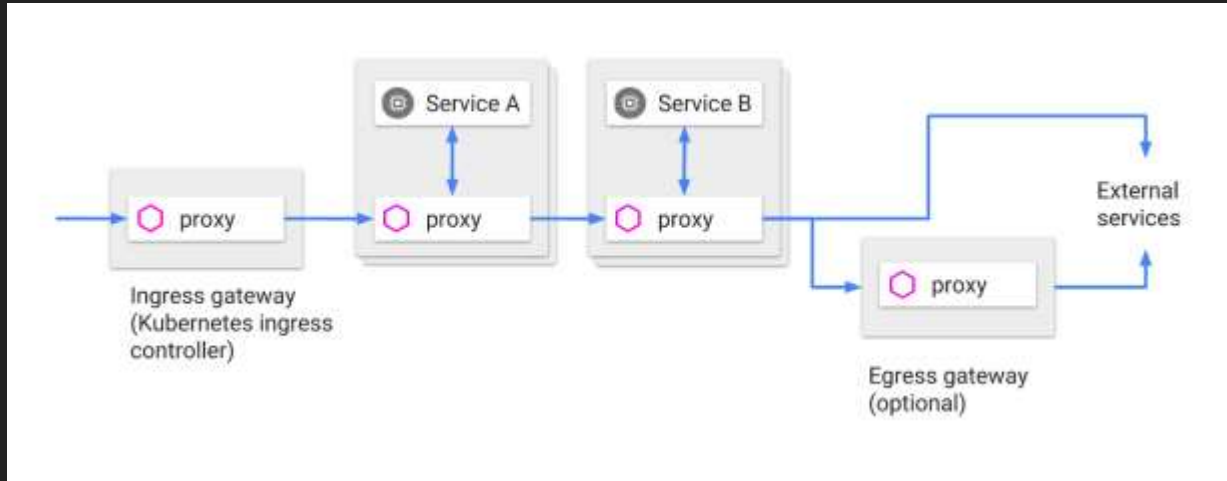
Cited from <https://www.infoq.com/presentations/istio-service-mesh>



<https://github.com/envoyproxy/envoy>

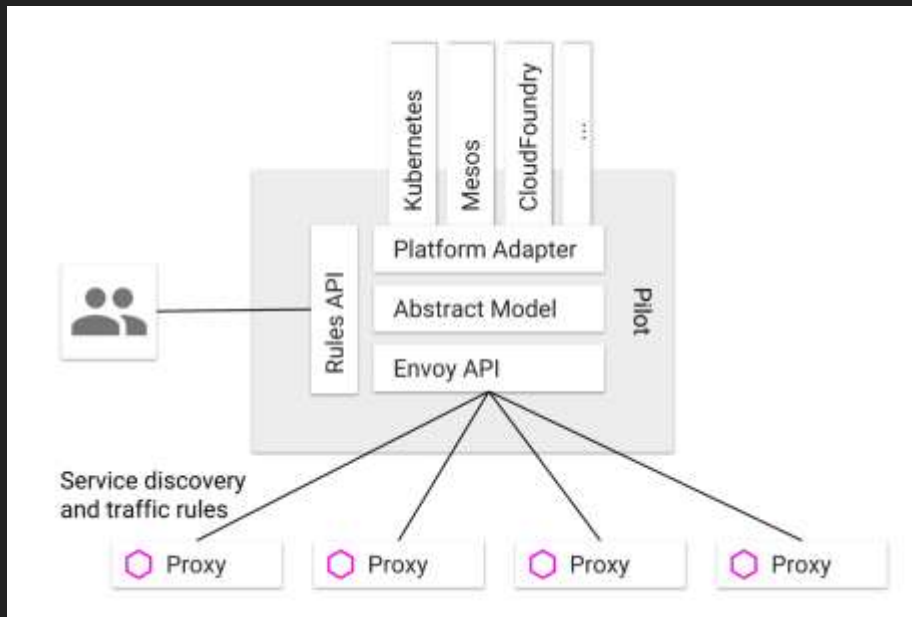
- L7 proxy and communication bus
- Written in C++11
- Pluggable architecture (L3 and L7)
- HTTP/1.1, HTTP/2 and gRPC
- Health checking
- L7 routing
- Advanced loadbalancer
- Dynamic configuration
- Metrics and tracing
- Battle tested at Lyft

ROUTING CONTROL



- **Ingress**
 - **Gateway**: configures a load balancer (one or multiple FQDN) for HTTP/TCP traffic at the edge of the mesh
- **Inside mesh**
 - **VirtualService**: defines the rules that control how requests for a service (one or multiple FQDN) are routed.
 - **DestinationRule**: configures the set of policies (a set of instances) to be applied to a request after VirtualService routing has occurred
- **Egress**
 - **ServiceEntry**: commonly used to enable requests to services outside of an Istio service mesh

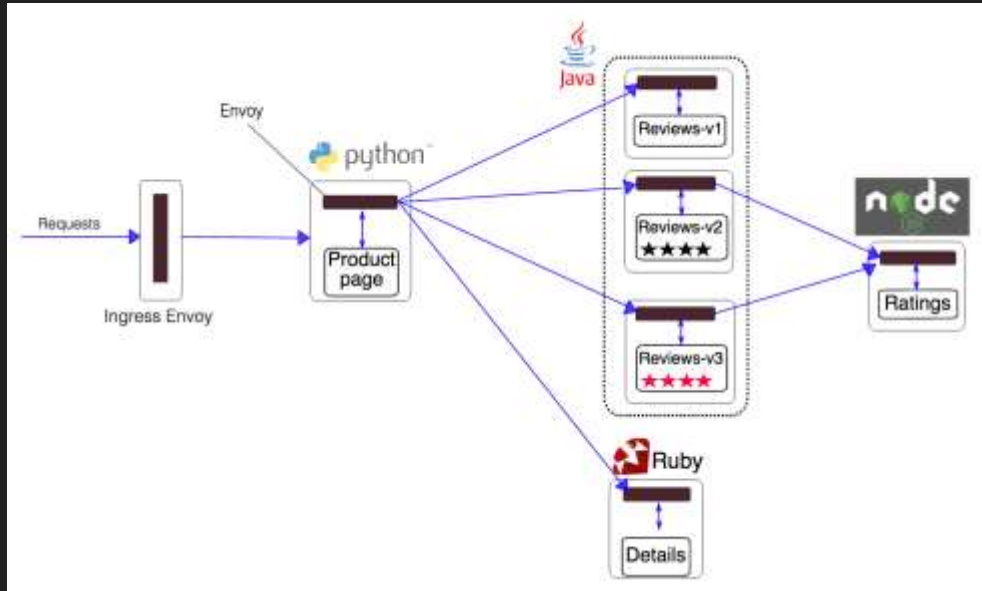
TRAFFIC MANAGEMENT - PILOT



- Maintain platform-agnostic model of services in the mesh
- Platform-specific adapter implement platform specific logic
 - Service discovery
 - Ingress resource
 - Rule definition
- Push configuration to Envoy and apply without restarts

TRAFFIC MANAGEMENT

REQUEST ROUTE (DEMO)

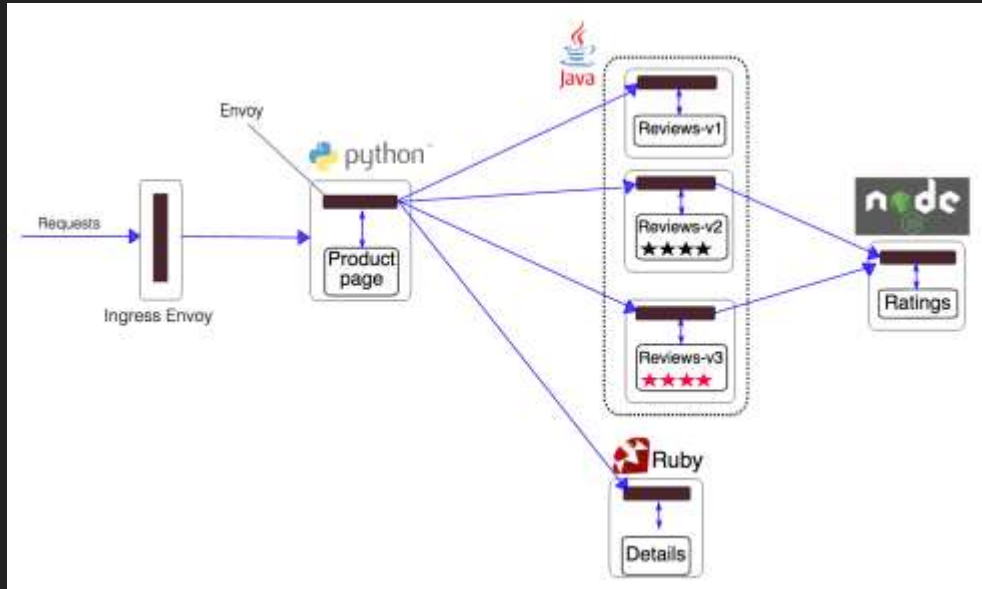


- Route to specific version of service
- Route based on request attribute
- ...

Three versions of the service of reviews

TRAFFIC MANAGEMENT

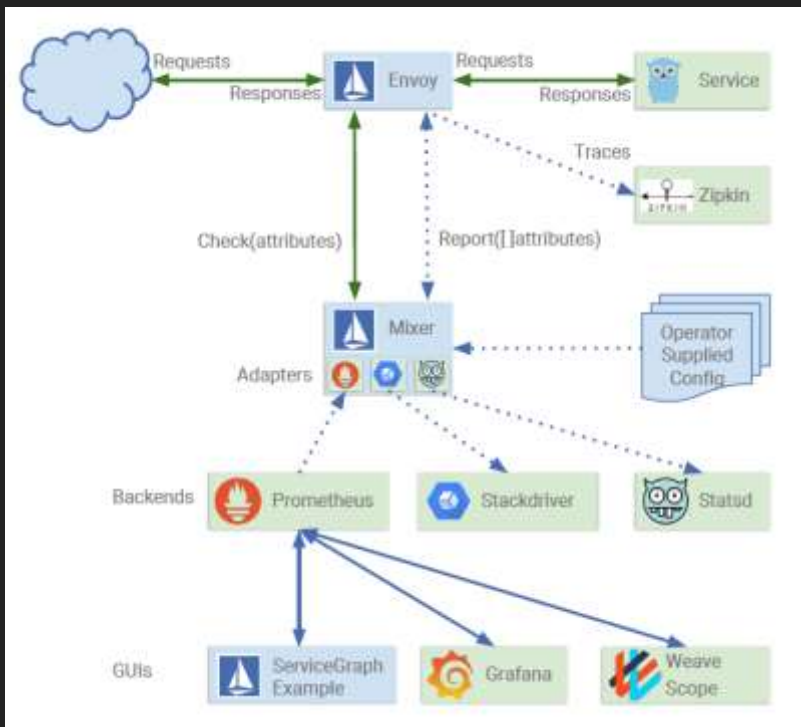
TRAFFIC SHIFT (DEMO)



Three versions of the service of reviews

- Route to existing version of service
- Route part of the request to new version of service
- Route all the request to the new version of service

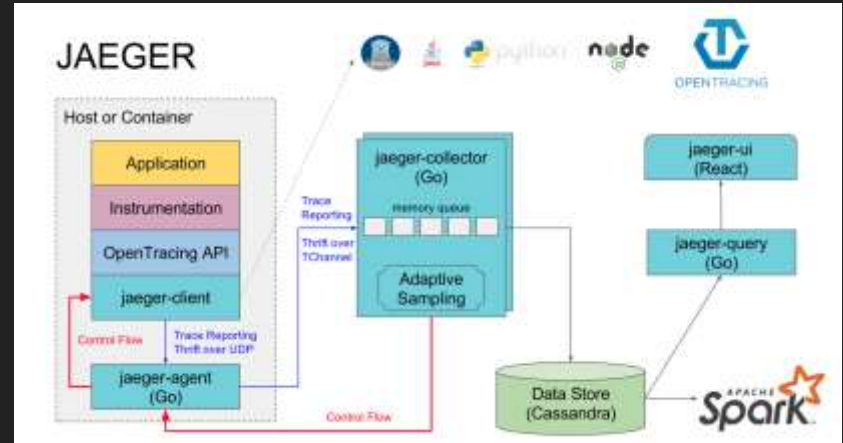
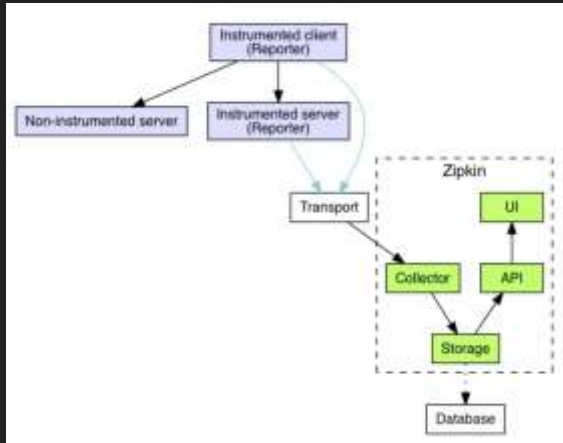
OBSERVABILITY - MIXER



- Collect metrics and logs emitted by Envoys without instrumenting apps
- Provide a uniform abstraction between application and infra backend
- Adapters in the Mixer normalize and forward to backends (monitoring, billing ...)
- Trace flow of requests across services
- Mixer is stateless with caching and buffering

OBSERVABILITY

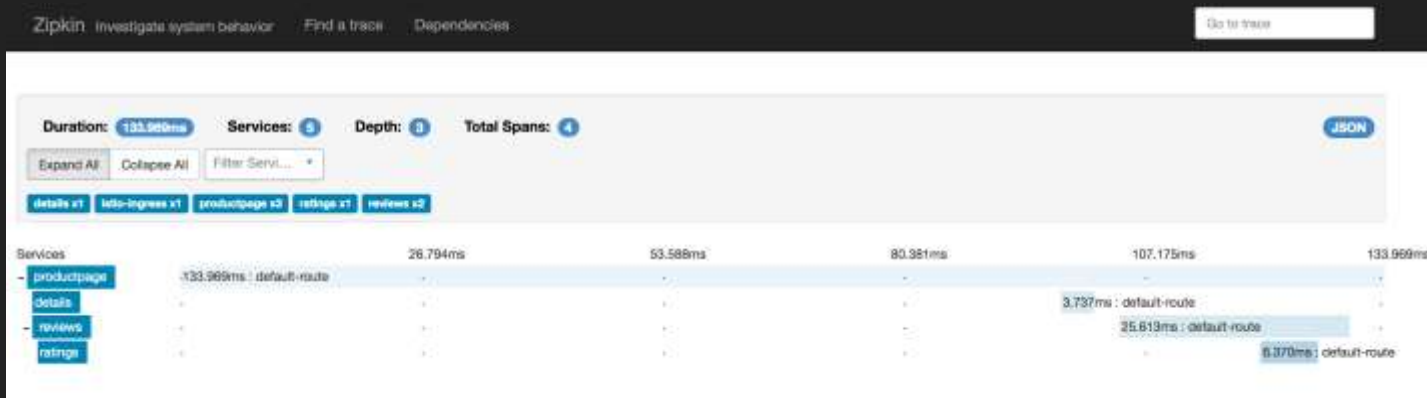
TRACING WITH ZIPKIN AND JAEGER



- Instrumentation
- Transport or agent
- Collector
- Storage
- API and UI

OBSERVABILITY

TRACING WITH ZIPKIN AND JAEGER (DEMO)



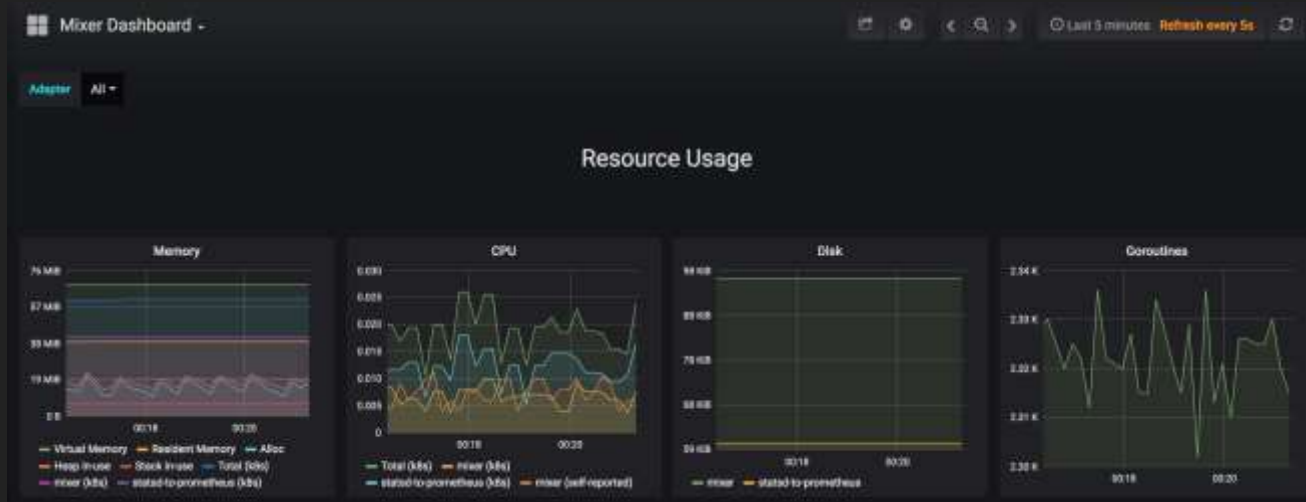
Application is required to collect and propagate the following headers from the incoming request to any outgoing requests

- x-request-id
- x-b3-traceid
- x-b3-spanid
- x-b3-parentspanid
- x-b3-sampled
- x-b3-flags
- x-ot-span-context

Zipkin has been replaced by Jaeger in Istio 0.8 and newer version

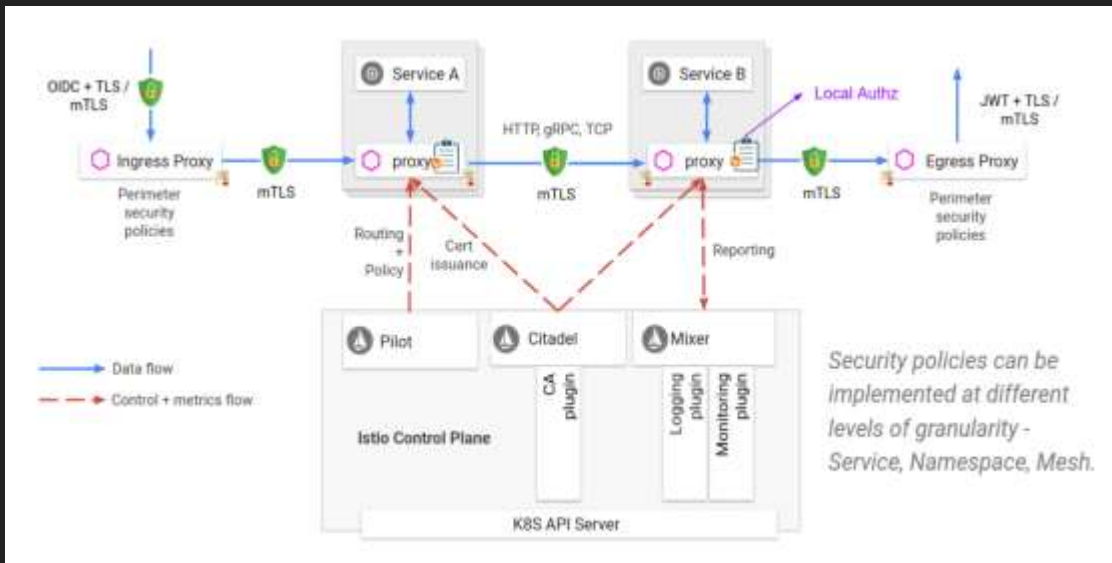
OBSERVABILITY

METRICS WITH MIXER + PROMETHEUS (DEMO)



- Generation of instances (in this example, metric values) from Istio attributes
- Creation of handlers (configured Mixer adapters) capable of processing generated instances
- Dispatch of instances to handlers according to a set of rules

SECURITY

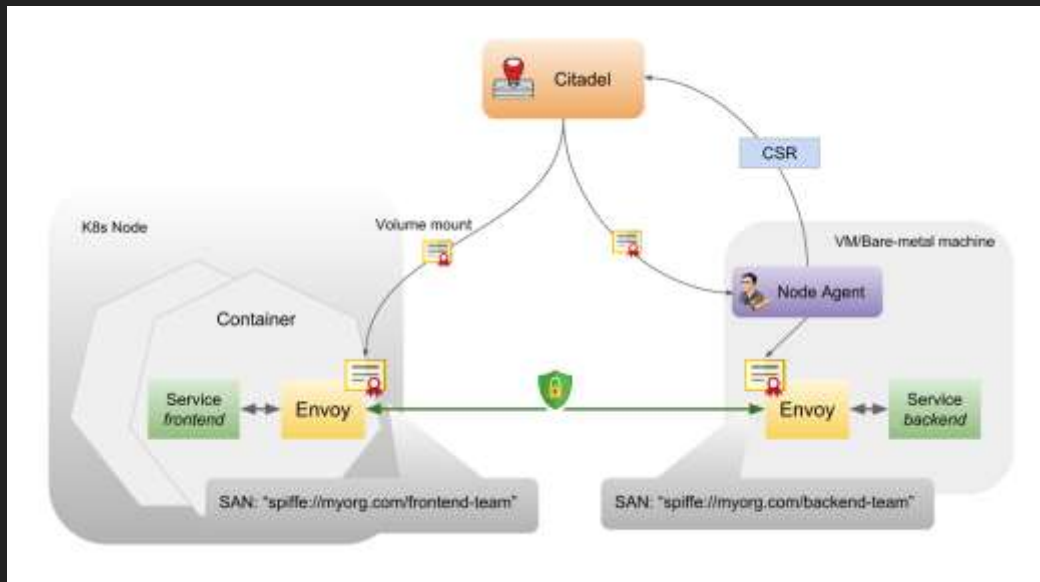


Need a management service to

- Secure microservices and their communication without instrumenting apps
 - ID management
 - Key/cert management
 - Access control

SECURITY

CERT MANAGEMENT - CITADEL

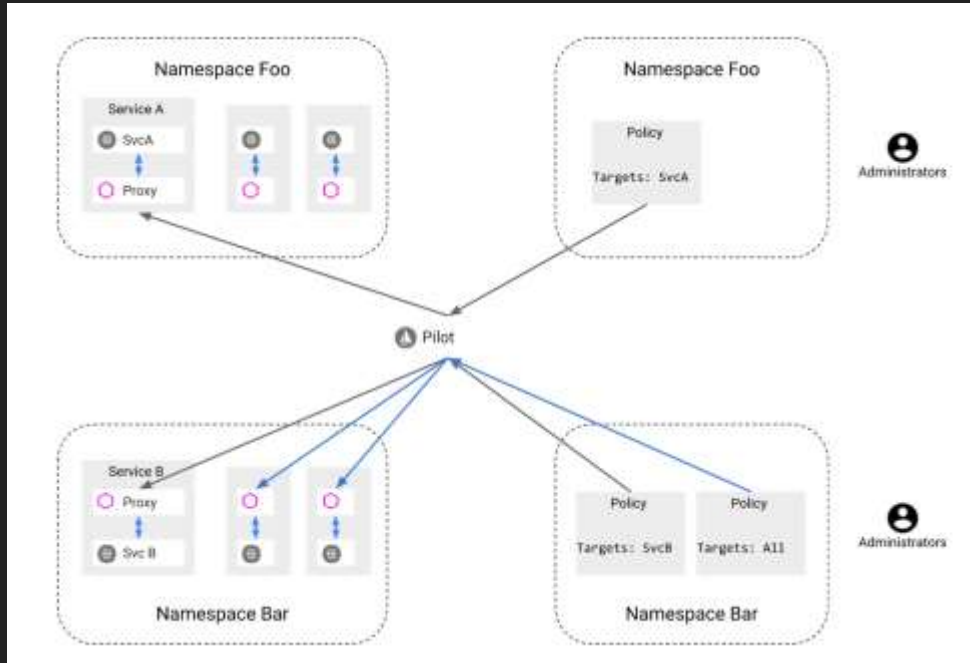


- Key management system
 - Automate key and certificate generation for service account
 - Distribute key/certs as kubernetes secret
 - Rotate keys/certs periodically
 - Revoke key/certs when required
 - Self-signed (root) CA vs user defined (root) CA
- Identity management
 - Identify service by service account in SPIFFE format

`"spiffe://<domain>/ns/<namespace>/sa/<serviceaccount>"`

SECURITY

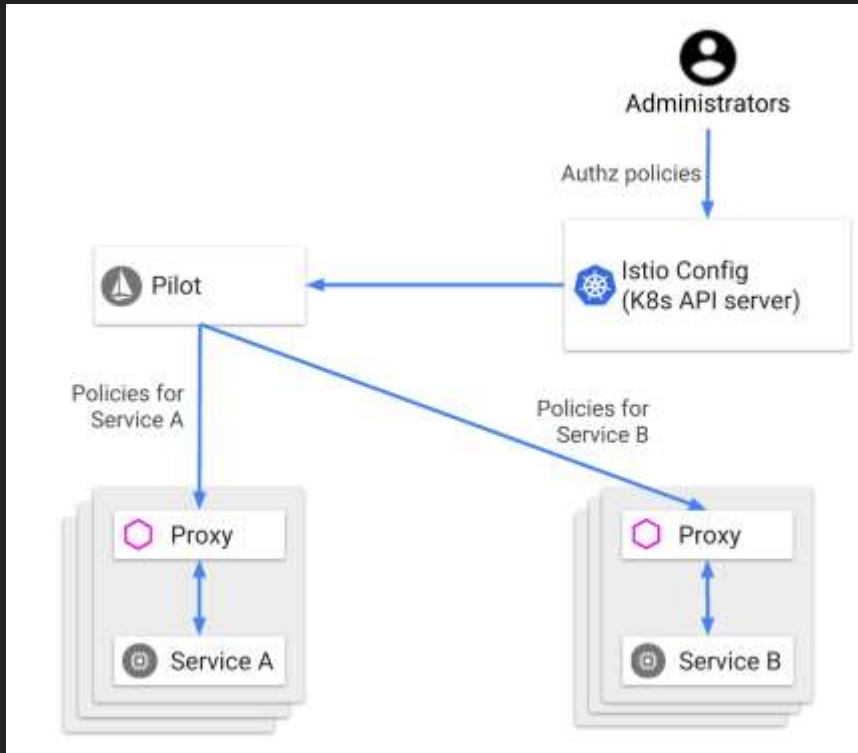
ID MANAGEMENT - PILOT



- Generate the config with proper information of key/certs and secure naming information, and then pass to envoy

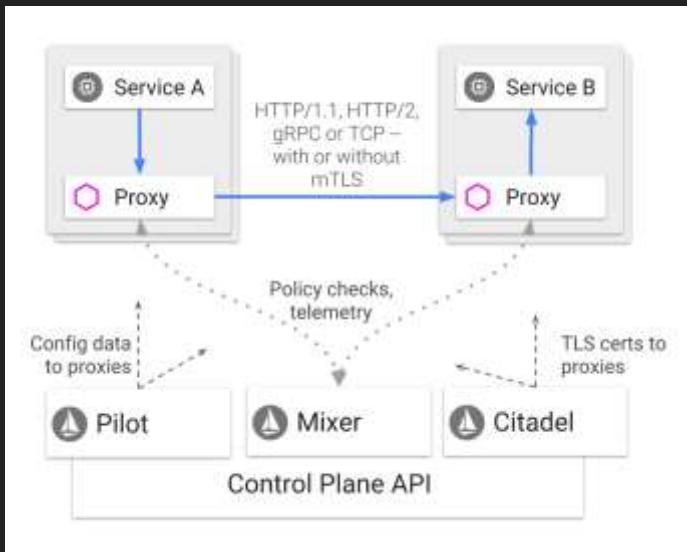
SECURITY

ACCESS CONTROL - PILOT



- Access control between services
 - Rule based by defining “Rule” with “match”, “handler” and “instance”
 - RBAC based by defining “ServiceRole” with “action”, and defining “ServiceRoleBinding” with “subject” and “roleRef”

KUBERNETES INTEGRATION



- Inject sidecar by leveraging MutatingAdmissionWebhook
- Manage policy and configuration by leveraging Custom Resource Definitions
- Identify account by leveraging Service Account and Secret
- Identify service by leveraging Service Discovery

PERFORMANCE AND SCALABILITY

- 1 vCPU per peak thousand requests per second for the sidecar(s) with access logging (which is on by default) and 0.5 without, fluentd on the node is a big contributor to that cost as it captures and uploads logs.
- Assuming typical cache hit ratio (>80%) for mixer checks: 0.5 vCPU per peak thousand requests per second for the mixer pods.
- Latency cost/overhead is approximately 10 millisecond for service-to-service (2 proxies involved, mixer telemetry and checks) as of 0.7.1, we expect to bring this down to a low single digit ms.
- mTLS costs are negligible on AES-NI (aes in /proc/cpuinfo) capable hardware in terms of both CPU and latency.

USEFUL LINKS

- ▶ Web istio.io
- ▶ Twitter: [@Istiomesh](https://twitter.com/Istiomesh)
- ▶ Github: <https://github.com/istio/istio>
- ▶ Community Doc: <https://istio.io/docs>
- ▶ Traffic management using Istio: <https://ibm.co/2F7xSnf>
- ▶ Resiliency and fault-tolerance using Istio:
<https://bit.ly/2qStF2B>
- ▶ Reliable application roll out and operations using Istio:
<https://bit.ly/2K9IRQX>

