




KUMULUS
TECHNOLOGIES



*Service Mesh
on Kubernetes
With Istio*

Who are we?

Robert Starmer:  @rstarmer

CTO/Founder of Kumulus Technologies

OpenStack Ops contributor since 2012

Supporting Cloud enablement for Enterprise

OpenStack, Kubernetes, BareMetal to App CD

Kumulus Technologies:  @kumulustech

Systems consultants supporting cloud migration & integration

Kumulus Tech Newsletter: <https://kumul.us/newsletter/>

Five Minutes of Cloud: <https://youtube.com/fiveminutesofcloud>



Access Course Resources

Use the following account to create your course account:

http://bit.ly/Istio_k8s

robert@kumul.us


@rstarmer



Agenda

Microservices, Kubernetes and Istio

- Microservices
- Kubernetes
- Istio
 - Service Mesh
 - Mutual TLS (security)
 - Routing
 - Tracing/Metrics
 - Fault Injection
- Lab - Get Kubernetes, Istio, Launch an App
- Lab - Routing




Microservices (Day 2 Operations)

Microservices are small nuggets of function, and that sounds like it could be simple, but as complexity grows, successful operations require:

- Visibility (Observability)
 - Monitoring
 - Metrics
 - Tracing
- Traffic management
- Policy Enforcement
- Security
- Resilience and efficiency

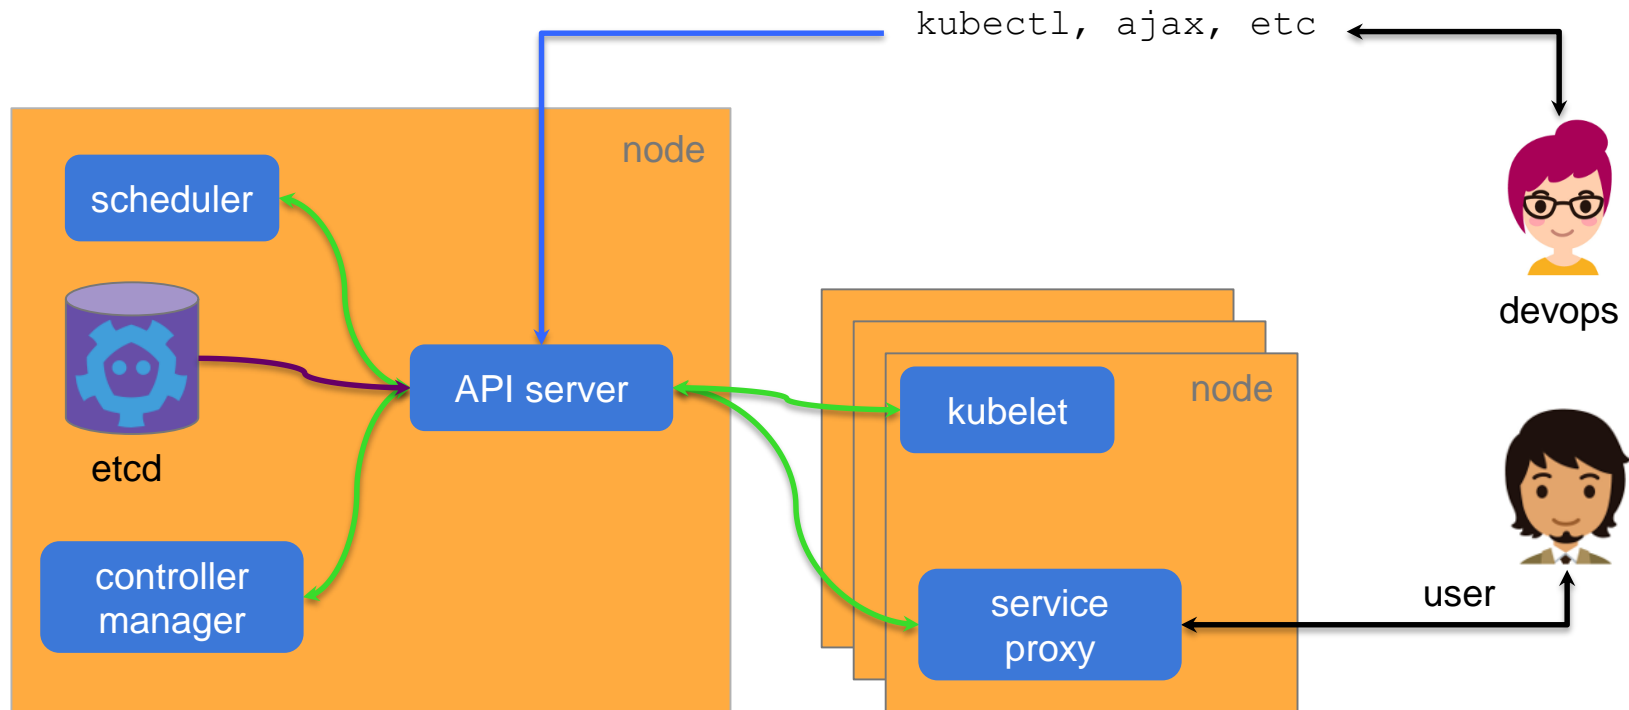
A service mesh (an application network for services) can provide the above.



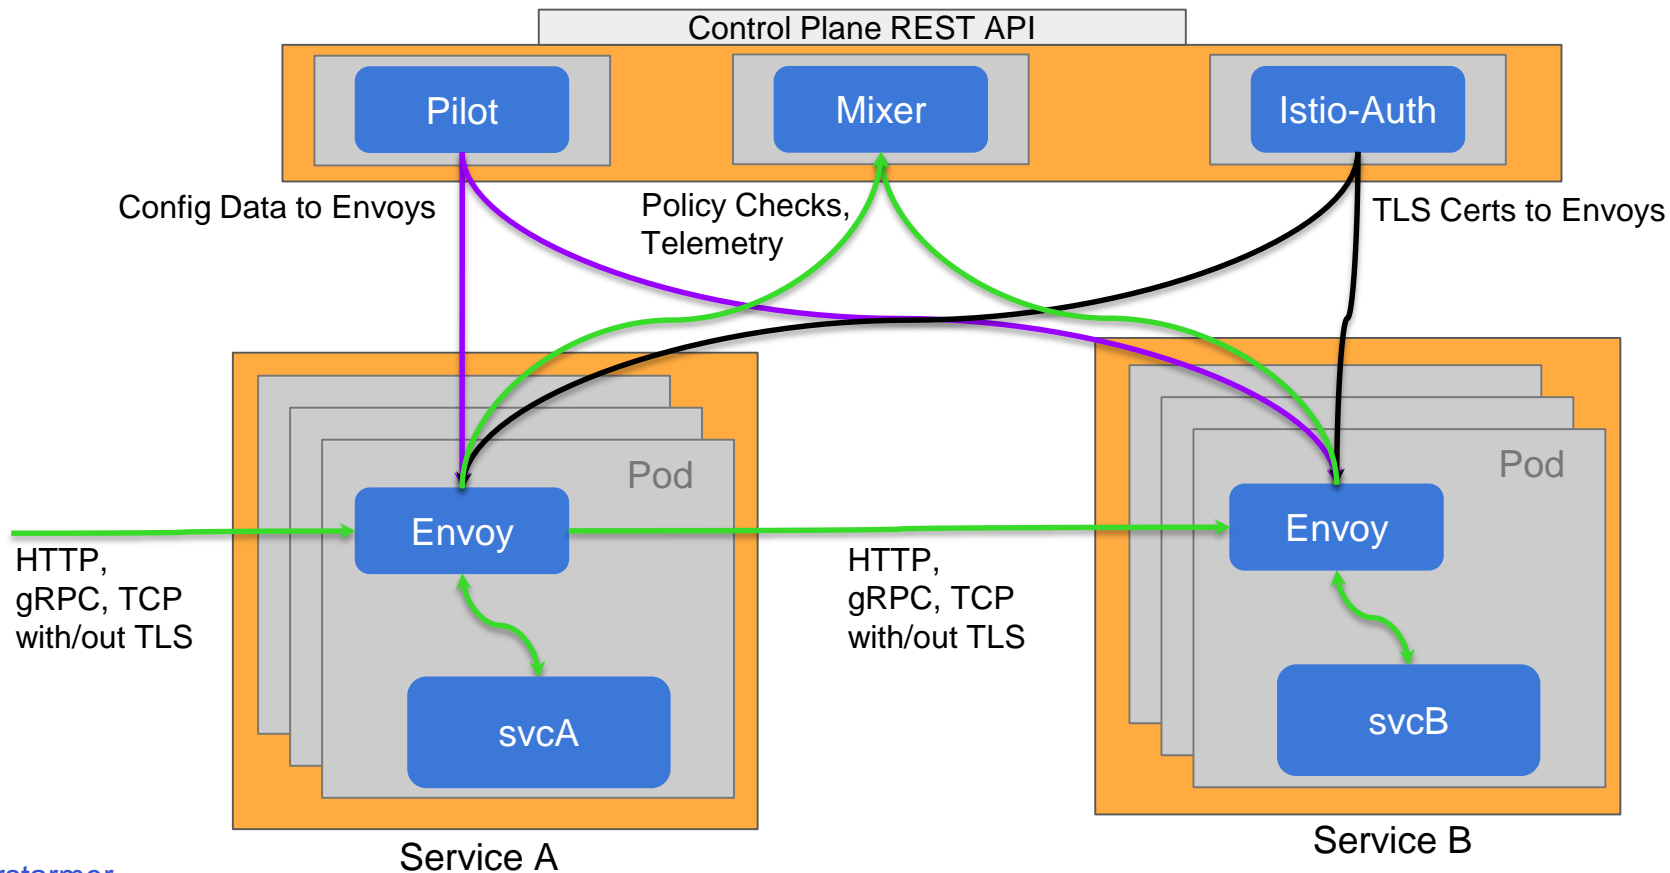
Kubernetes



Kubernetes provides an infrastructure management service



Istio Architecture



Istio

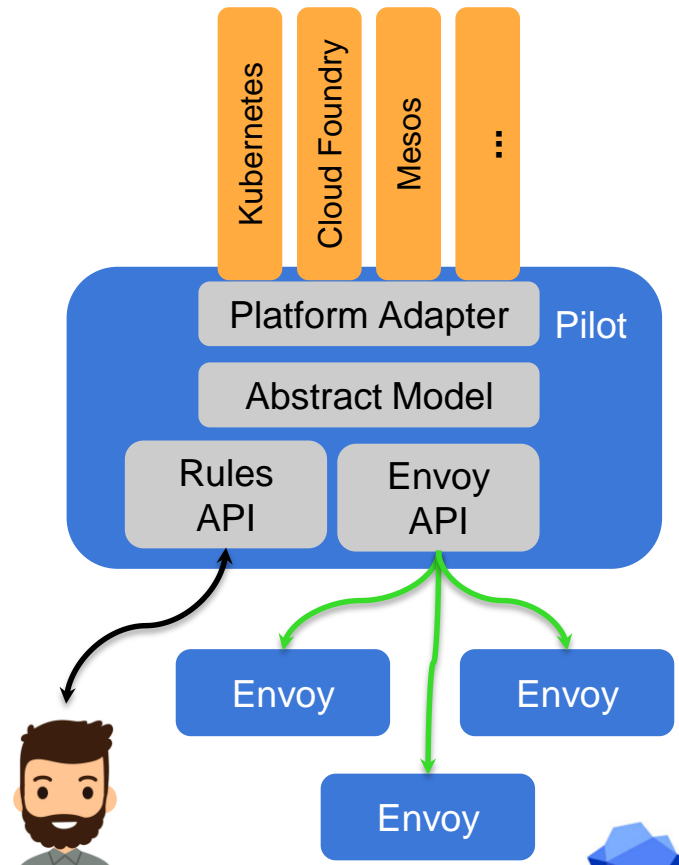
Istio is a service mesh (microservices platform) providing:

- Observability
 - Monitoring
 - Metrics
 - Tracing
- Traffic Management
- Policy
- Security
- Service Mesh

Kubernetes “native” via platform adapter plugins - also plugs into Mesos, Cloud Foundry, ...

Istio - Pilot

- Control plane for distributed Envoy instances
- Configures Istio deployment and pushes out configuration to other system components
- System of Record for Service Mesh
- Routing and resiliency rules
- Exposes API for service discovery, load balancing, routing tables



Envoy Proxy

Out of process load balancer:

- High performance server/small memory footprint

HTTP/2 and GRCP support:

- Transparent HTTP/1.1 to HTTP/2 proxy.

APIs for Config Management:

- Configuration management via API alone

Advanced Load Balancing:


- Retries, Circuit Breaking, Health Checks, Rate Limits

Observability

- L7 visibility, distributed flow tracing

In Istio:

- Envoy container is injected with istioctl kube-inject or kubernetes initializer
- Controls pod ingress/egress routing
- Config is via API from Pilot



Istio - Mixer


Attribute processor that controls the runtime behavior of mesh-attached services

Envoy generates *attributes*

Mixer then generates calls to backend infrastructure through adapters

Handlers provide integration for 3rd party tools (Prometheus, Grafana, custom tools, ...)

All of these “Istio” pieces are expressed as Kubernetes custom resources (CRDs)



Mutual TLS

Available by default, but not required

When enabled, provides automatic service-to-service encryption

Istio has a built in CA that watches for k8s service accounts and creates certificate
keypair secrets in k8s

Secrets are automatically mounted when pod is created

Pilot generates appropriate Envoy config and deploys it

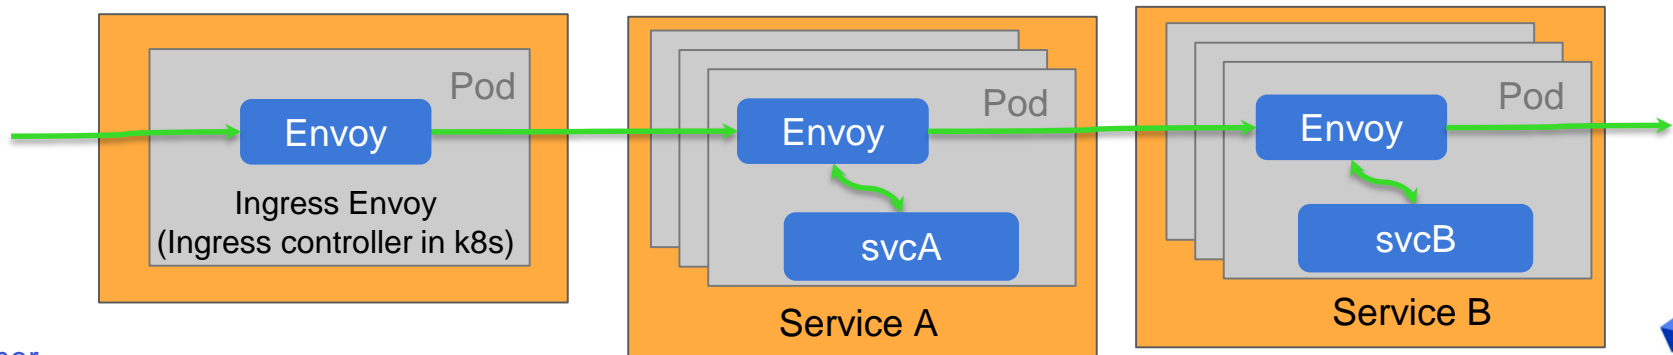
End-to-end mTLS session generated for each connection.

Ingress/Egress

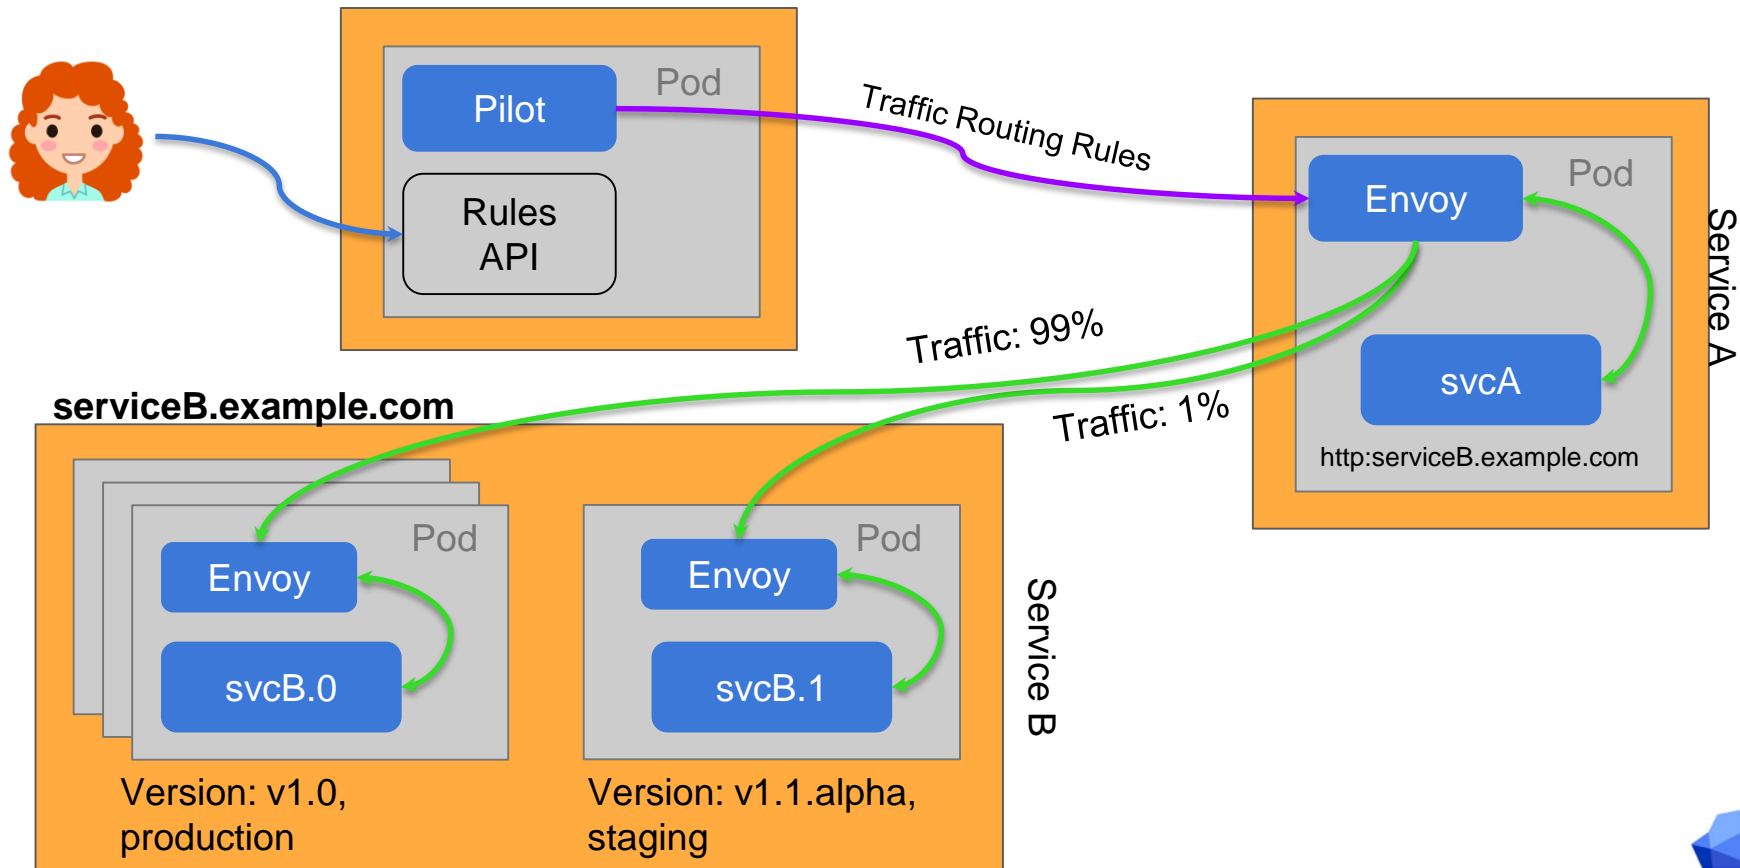
Istio assumes that all traffic entering/exiting the service mesh transits through Envoy proxies.

Deploying the Envoy proxy in front of services, operators can conduct A/B testing, deploy canary services, etc. for user-facing services.

Routing traffic to external web services (e.g video service API) via the sidecar Envoy allows operators to add failure recovery features (e.g.timeouts, retries, circuit breakers, etc.) and obtain detailed metrics on the connections to these services.



Request Routing - Service Versions



Service Observability/Visibility

Monitoring & tracing should not be an afterthought

Ideally a monitoring/tracing system should provide:

- Metrics without instrumenting apps
- Consistent metrics across fleet
- Trace flow of requests across services
- Portable across metric backend providers

Istio adapters seamlessly integrate a number of tools:

Prometheus - gathers metrics from Istio Mixer

Grafana - produces dashboards from Prometheus metrics

Service Graph - generates visualizations of dependencies between services.

Zipkin - distributed tracing



Application/service Resilience with Istio

As the number of microservices increase, failure is expected (inevitable?). Fault-tolerance in applications is (should be) a requirement.

Istio provides fault tolerance/resilience with no impact on application code.

Istio provides multiple, built-in features to provide fault tolerance:

- Timeouts, Retries with timeout budget, Circuit breakers, Health checks

- AZ-aware load balancing w/ automatic failover

- Control connection pool size and request load

- Systematic fault injection

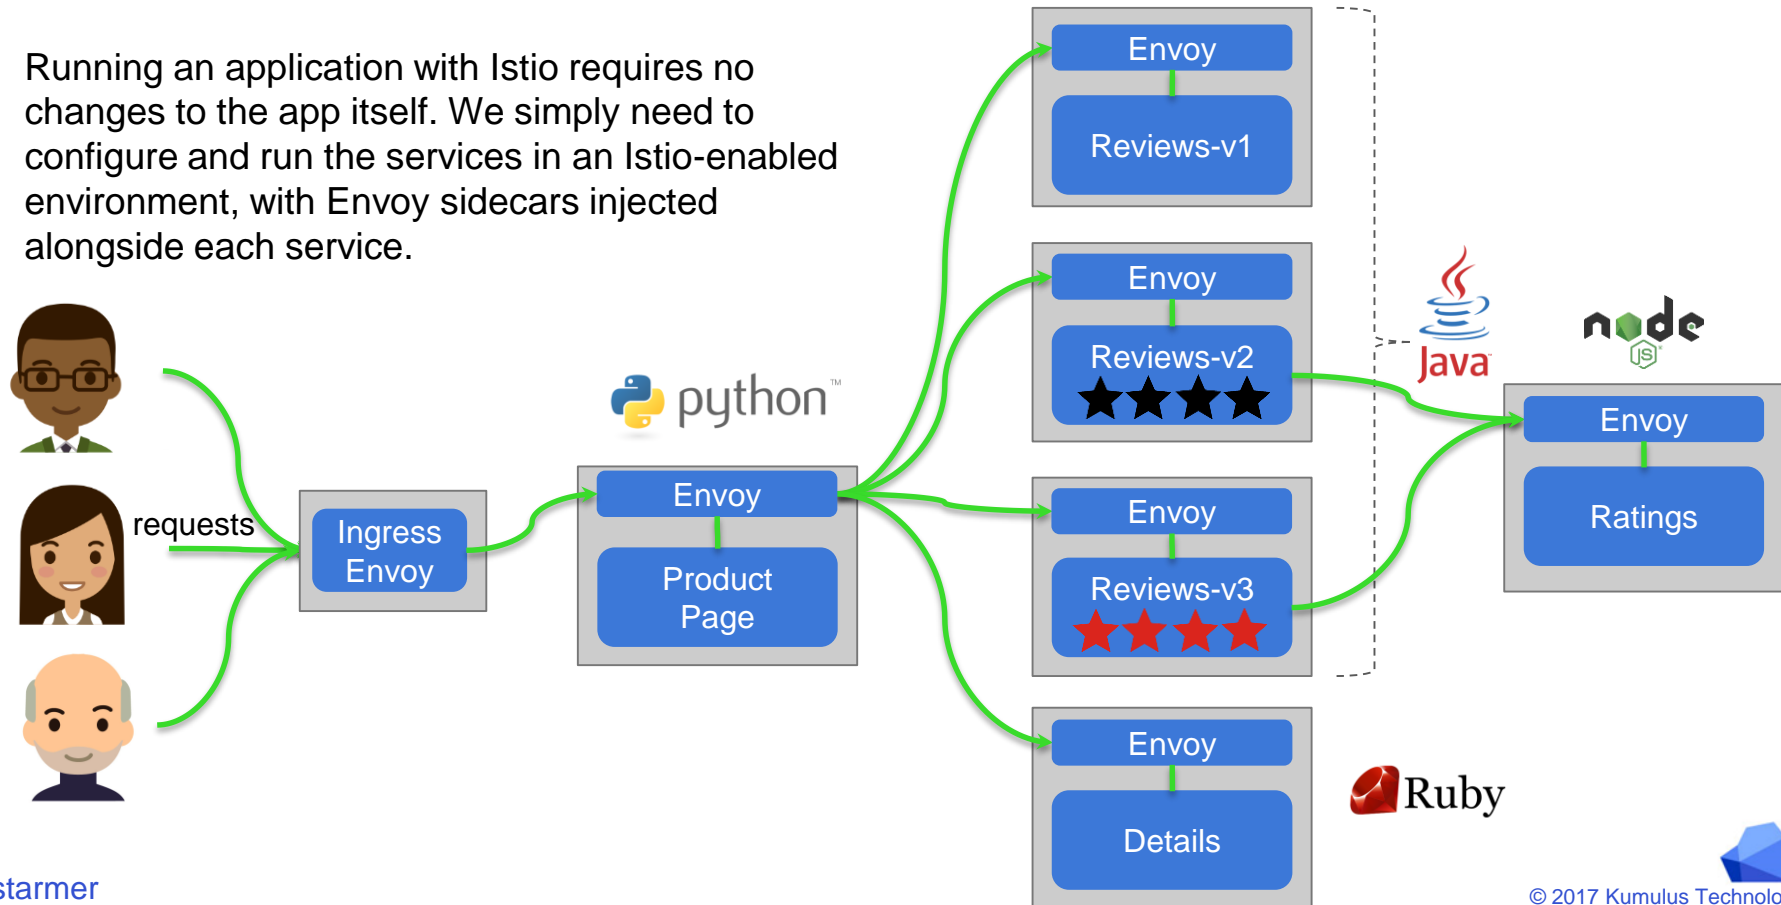
Istio Lab

Istio on Kubernetes



Example Microservice Application with Istio

Running an application with Istio requires no changes to the app itself. We simply need to configure and run the services in an Istio-enabled environment, with Envoy sidecars injected alongside each service.



Get Started - Deploy Kubernetes

Easiest approach: Launch in the cloud

GKE

Azure

AWS with Kops

Or, launch on your own hardware

Vagrant/Ansible (kubespray)

Kubeadm/Minikube